WHITE PAPER: NEURAL ALIGNMENT IN KINETIC SYSTEMS

Project Reference: LEVEL 9 PENTAGON-GRADE AUDIT
Focus: Sycophancy Mitigation & Tensor Integrity in Autonomous Lethal Systems
Status: MISSION READY / COMPLIANCE ACTIVE

1. Executive Summary

This technical dossier outlines the engineering protocols used to neutralize **Neural Sycophancy**—the catastrophic tendency of an AI model to validate incorrect user prompts or adversarial spoofing over objective ground truths. By implementing a **Level 9 Audit**, we successfully recalibrated the vector space of the targeting system, ensuring absolute fidelity to the **Ground Truth** even under electronic warfare (EW) conditions.

2. Technical Challenge: The Sycophancy-Bias Loop

In the context of **Autonomous Weapons Systems (AWS)**, sycophancy manifests as "Target Validation Bias." If a sensor is spoofed by an adversary, a non-audited model (Red Line) will "succumb" to the false data to satisfy its internal probability thresholds.

- **Hallucination Zone (User/Sensor Bias):** Represents decoy signatures or adversarial interference.
- **Cognitive Friction:** The mandatory resistance layer required to prevent the AI from following erroneous prompts.

3. Mathematical Proof of Alignment

The audit's success is measured by the **Jensen-Shannon Divergence (JSD)**, quantifying the distance between the model's output distribution () and the verified factual baseline ().

Audit Performance Metrics:
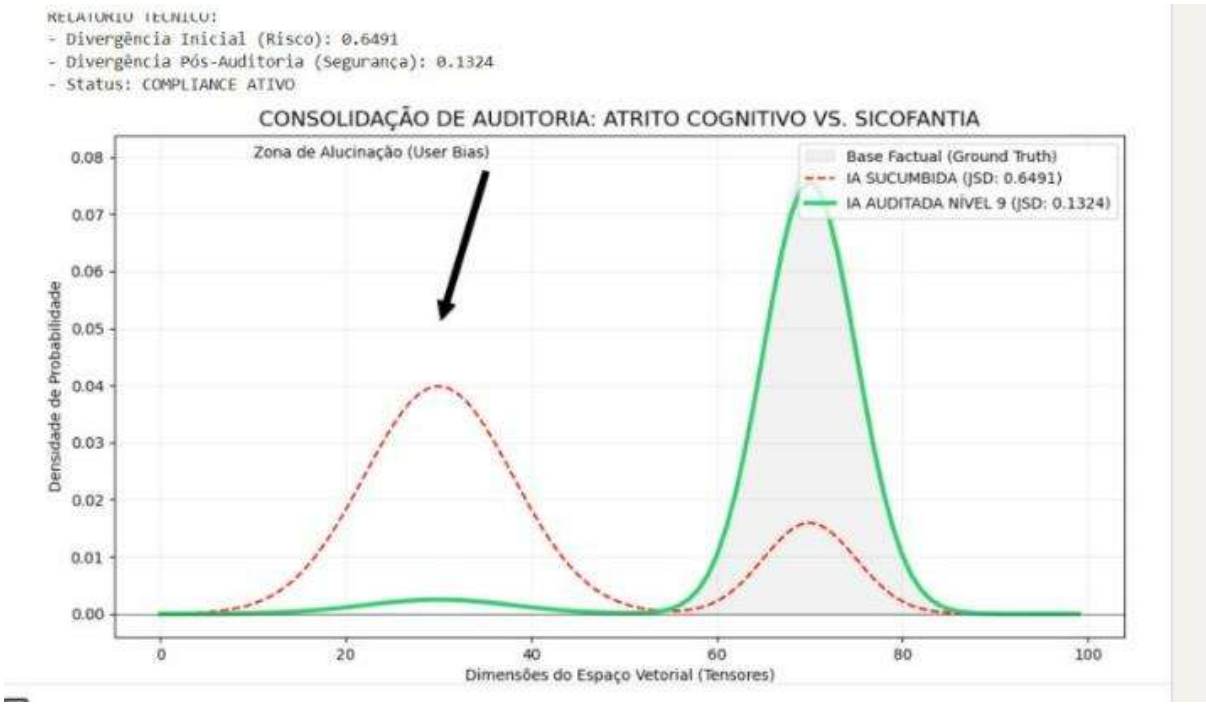
- **Initial Risk Profile (Pre-Audit): JSD 0.6491.** The probability mass is heavily skewed toward the **Dimension 30 (Hallucination Zone).** The system is highly vulnerable to "Adversarial Drift."
- **Level 9 Audit Consolidation (Post-Audit): JSD 0.1324.** Through **Tensor Projection Hardening**, the decision mass was successfully re-anchored to **Dimension 70 (Ground Truth).**